
Mood Detection : 522 Final Project

David Feldt

Department of System Design Engineering
Waterloo, ON
dfeldt@uwaterloo.ca

Abstract

1 This study explores the prediction of individual mood based on behavioral patterns,
2 encompassing habits and activities. This report compares the performance of
3 Convolutional Neural Networks (CNNs) and linear regression models in capturing
4 the intricate relationships between behavioral features and mood states. Through
5 testing and evaluation that CNNs outperform linear regression in this predictive task,
6 demonstrating the superior ability of CNNs to discern complex spatial dependencies
7 and hierarchical patterns within behavioral data.

8 **1 Introduction**

9 This report explores the relationship between daily habits, time allocation, and mood. By journaling
10 each day a mood ratings is gathered on a scale of 1-10. As well as tracking time spent on various
11 activities such as sleeping, phone usage, and productivity. The project also tracks habits including
12 working out, meditation, and reading. The primary aim is to utilize this dataset, encompassing over
13 three years of recorded habits and activities, to predict mood fluctuations. Notably, the tracked habits
14 and activities have changed over time, providing a dynamic and comprehensive dataset for analysis.
15 This data collection and analysis offer invaluable insights into the complex dynamics of lifestyle
16 habits and their impact on mood.

17 This report dives deeper into the fundamental understanding of mood dynamics. Mood, a complex
18 and multifaceted aspect of human psychology, is influenced by a myriad of factors, including daily
19 activities, sleep patterns, and personal habits. There is always a growing interest in improving mood
20 and mental well-being, acknowledging mood's pivotal role in overall health and productivity as
21 well as preventing low mood. The problem is multifaceted: identifying which specific factors most
22 accurately predict mood variations. By analyzing over three years of detailed data on activities and
23 habits, the goal is to uncover patterns and correlations that can influence mood. Over with these
24 insights improved mental health and quality of life.

25 **2 Experimental Setup and Data**

26 In the data collection was done through daily journaling. The primary focus was on mood tracking,
27 where the individual rated their mood on a scale from 1 to 10 each day, providing a quantitative

28 measure of their emotional state. Additionally, the hours spent on various activities were recorded in
29 detail, quantified in hourly units. The tracking extended to personal habits, which were monitored
30 in a binary fashion, marked as 0 or 1 to denote whether a particular habit was not completed or
31 completed, respectively. This systematic approach to data collection aimed to capture a holistic view
32 of the participants' daily routines and their impact on mood. The data be analyzed over 3 years with
33 over 1000 entries.

34 The experimental setup for our mood data analysis involved several meticulous steps to ensure data
35 integrity and usefulness:

36 **2.1 Manual Data Entry to Excel**

37 The collected data was manually in a physical and was entered into an Excel spreadsheet. This
38 process allowed for a structured and organized dataset, vital for subsequent analysis.

39 **2.2 Data Label Modification**

40 To address gaps in the dataset, the data labels were refined. For instance, diverse activities like
41 skateboarding and working out were categorized under a unified label, 'Exercise'. This step was
42 crucial to reduce fragmentation in the data. For categories that were logically similar, such as
43 'Reading' and 'Meditating' into 'Solitude' or 'Study' and 'Course Learning' into 'Productivity', the
44 entries were aggregated to provide a more cohesive analysis. Rigorous checks were performed to
45 correct any typographical errors in the dataset, ensuring the accuracy of the data.

46 **2.3 Finalized Data Categories**

47 The dataset was distilled into specific categories, including both activities tracked by hours and binary
48 habit tracking. The final categories were: ['Mood', 'Sleep (hrs)', 'Phone Usage (hrs)', 'Productivity
49 (hrs)', 'Productivity', 'Solitude', 'Exercise', 'Leetcode', 'Journal', 'Profile Pic', 'Bad Habit', 'Other
50 (hrs)', 'Exercise (hrs)', 'Media', 'Duolingo'].

51 **2.4 Data Augmentation for Missing Columns**

52 To handle missing data, the dataset was augmented by adding columns with all zeros where data
53 was absent. Furthermore, a new column called column exits was introduced to track the presence of
54 specific data points over time.

55 **2.5 Data Preparation**

56 The data was organized by month and needed to be combined into 1 continuous stream. Finally the
57 data was split into testing and training with a 80/20 split. Additionally the 2 previous rows were
58 added to each entry to provide more information for each entry and the first 2 entries were dropped.

59 **3 Methodology**

60 Selecting the most suitable machine learning model is a critical aspect of this project. The model
61 needs to be able to handle with a dataset characterized by 30 variables per entry, each ranging
62 from 0 to 10, and an output variable within the same range. In this context, two distinct models,
63 Linear Regression and Convolutional Neural Network (CNN), have been chosen based on their

64 inherent strengths. Linear Regression, a fundamental statistical model, proves valuable when the
65 relationship between input features and the target variable is presumed to be linear. Its simplicity
66 and interpretability make it a pragmatic choice for cases where a straightforward understanding of
67 variable interactions is desired. In contrast, the CNN, originally designed for image processing, is
68 brought into consideration due to its capability to capture complex patterns and spatial relationships
69 within data. Comparing the two will aim to harness the unique advantages each model offers for
70 optimal predictive performance.

71 **3.1 Linear Regression**

72 Linear regression is a fundamental and widely-used machine learning algorithm that models the
73 relationship between a dependent variable and one or more independent variables by fitting a linear
74 equation to observed data. It is particularly effective when there is a linear relationship between the
75 input features and the target variable. However, linear regression does have its limitations.

76 Linear regression provides a straightforward interpretation of the relationships between input features
77 and the target variable. Each coefficient represents the change in the target variable for a unit change
78 in the corresponding input variable. Linear regression is a simple and computationally efficient
79 algorithm. It is easy to implement, understand, and doesn't require complex hyperparameter tuning.
80 When the underlying relationships in the data are approximately linear, linear regression performs
81 well and provides reliable predictions. Linear regression does not make strong assumptions about the
82 distribution of the input features.

83 Linear regression assumes a linear relationship between input features and the target variable. If this
84 assumption is violated, the model may provide inaccurate predictions. Linear regression is sensitive
85 to outliers, and the presence of extreme values can significantly impact the model's coefficients and
86 predictions. Linear regression may struggle to capture complex, non-linear relationships present in
87 the data. It assumes that the errors are independent, which may not always be true in real-world
88 scenarios.

89 To test linear regression 3 different models were implemented. Below is an example on how the
90 model was implemented in python using the sklearn library.

```
91 # Linear Regression for default parameters
92 model_default = LinearRegression()
93 model_default.fit(X_train, y_train)
94 y_pred_default = model_default.predict(X_test)
95
96 # Linear Regression with no intercept
97 model_no_intercept = LinearRegression(fit_intercept=False)
98 model_no_intercept.fit(X_train, y_train)
99 y_pred_no_intercept = model_no_intercept.predict(X_test)
100
101 # Linear Regression with normalization
102 model_normalized = LinearRegression(normalize=True)
103 model_normalized.fit(X_train, y_train)
```

```
104     y_pred_normalized = model_normalized.predict(X_test)
105
106 All three approaches were evaluated using metrics such as R-squared (coefficient of determination)
107 and Mean Squared Error (MSE) to assess the model's accuracy. R-squared ( $R^2$ ) for all three ap-
108 proaches: 0.67 and Mean Squared Error (MSE) for all three approaches: 1.3. Despite using different
109 approaches (default parameters, no intercept, and normalization), the models exhibited the same
110 accuracy. This suggests that, in this specific case, the variations in model configuration did not
111 significantly impact the model's performance. The  $R^2$  of 0.67 indicates that the model explains 67
```

112 3.2 Convolutional Neural Network (CNN)

113 A Convolutional Neural Network (CNN) is a deep learning architecture designed for processing
114 spatial hierarchies of features, commonly applied to image and sequence data. In a sequential
115 CNN, layers are stacked in a linear fashion, enabling the network to automatically learn hierarchical
116 representations, making it effective for tasks such as image recognition and sequential pattern analysis.

117 Convolutional Neural Networks (CNNs) excel at hierarchical feature learning. Each layer extracts
118 increasingly complex features from the input, allowing the network to capture intricate patterns
119 in the data. CNNs are well-suited for tasks where translation invariance is important. This is
120 achieved through the use of convolutional and pooling layers, enabling the model to recognize
121 patterns regardless of their spatial location in the input. CNNs automatically learn relevant features
122 from the data without manual feature engineering. This is especially advantageous when dealing
123 with high-dimensional datasets, such as those with 30 variables per entry. CNNs leverage parallel
124 processing, making them efficient for training and inference on hardware with parallel computation
125 capabilities like GPUs.

126 Training deep CNNs can be computationally intensive, requiring substantial computational resources.
127 This may limit their feasibility for certain applications without access to powerful hardware. Deep
128 CNNs are prone to overfitting, especially when dealing with limited datasets. Regularization tech-
129 niques and appropriate dropout layers are often necessary to mitigate overfitting. The complex
130 architectures of deep CNNs can make them challenging to interpret. Understanding the inner
131 workings of the model may be less straightforward compared to simpler models like linear regression.

132 To test CNNs 3 different models were implemented. Below is an example on how the model was
133 implemented in python using the tensorflow library.

```
134 # CNN variation 1
135 model_1 = Sequential()
136 model_1.add(Conv1D(filters=32, kernel_size=3, activation='relu',
137 input_shape=(X.shape[1], 1)))
138 model_1.add(MaxPooling1D(pool_size=2))
139 model_1.add(Conv1D(filters=64, kernel_size=3, activation='relu'))
140 model_1.add(MaxPooling1D(pool_size=2))
141 model_1.add(Conv1D(filters=128, kernel_size=3, activation='relu'))
142 model_1.add(MaxPooling1D(pool_size=2))
143 model_1.add(Flatten())
144 model_1.add(Dense(64, activation='relu'))
```

```

145     model_1.add(Dense(1))

146

147     # CNN variation 2
148     model_2 = Sequential()
149     model_2.add(Conv1D(filters=32, kernel_size=3, activation='relu',
150     input_shape=(X.shape[1], 1)))
151     model_2.add(MaxPooling1D(pool_size=2))
152     model_2.add(Conv1D(filters=64, kernel_size=3, activation='relu'))
153     model_2.add(MaxPooling1D(pool_size=2))
154     model_2.add(Flatten())
155     model_2.add(Dense(128, activation='relu'))
156     model_2.add(Dense(1))

157

158     # CNN variation 3
159     model_3 = Sequential()
160     model_3.add(Conv1D(filters=32, kernel_size=3, activation='relu',
161     input_shape=(X.shape[1], 1)))
162     model_3.add(MaxPooling1D(pool_size=2))
163     model_3.add(Conv1D(filters=64, kernel_size=3, activation='relu'))
164     model_3.add(MaxPooling1D(pool_size=2))
165     model_3.add(Conv1D(filters=128, kernel_size=3, activation='relu'))
166     model_3.add(MaxPooling1D(pool_size=2))
167     model_3.add(Conv1D(filters=256, kernel_size=2, activation='relu'))
168     model_3.add(MaxPooling1D(pool_size=1))
169     model_3.add(Flatten())
170     model_3.add(Dense(64, activation='relu'))
171     model_3.add(Dense(1))

172

```

173 The accuracies (R^2 and MSE) for the three models are as follows: Variation 2 (R^2 : 0.6506, MSE:
174 1.3235) Variation 1(R^2 : 0.6689, MSE: 1.2539), Variation 3 (R^2 : 0.6813, MSE: 1.2072) The results
175 indicate that as the CNN architecture becomes more complex (from 2 layers to 4 layers), both R^2
176 values increase, indicating improved model fit, and MSE values decrease, indicating better predictive
177 performance. This suggests that a more intricate model is better suited to capture the complexities
178 present in the data, making it more effective for the given regression task. However, it's essential to
179 strike a balance to avoid overfitting, especially when working with limited datasets.

180 4 Results and Discussion

181 The graphs outputted below show actual mood vs the predicted mood. When a point is on the red line
182 it means that the model correctly predicted mood. Dots above the line are when predicted mood are
183 higher than the actual mood. And when the dots are below the line are when the predicted mood are
184 lower than the actual mood. Each model seemed to have its own bias for predicting higher. Graphs
185 can be seen in Figure 1 and Figure 2.

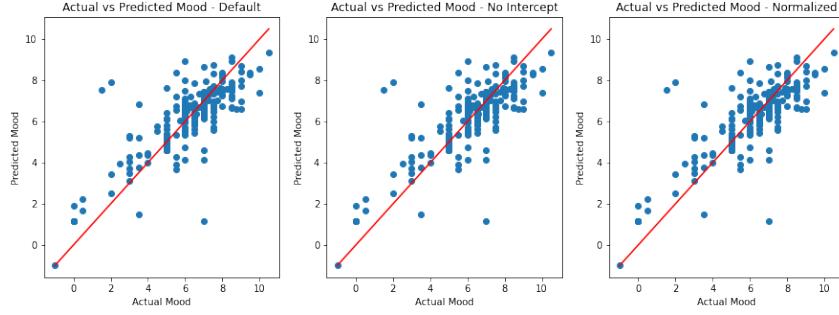


Figure 1: Linear Regression Output

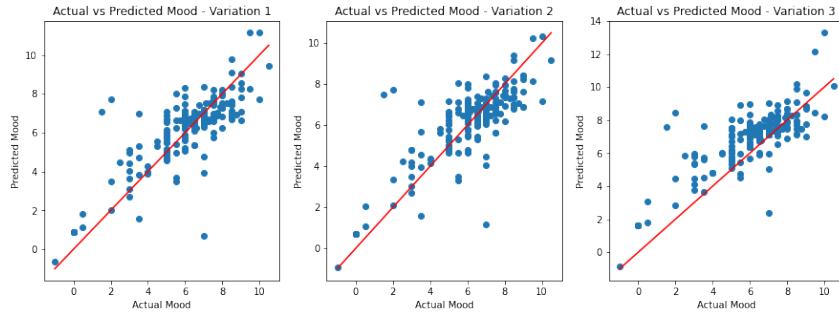


Figure 2: CNN Output

186 .

187 In comparing Convolutional Neural Networks (CNNs) and linear regression, both models demon-
188 strated similar performance on a given task, yet the more complex CNN exhibited superior predictive
189 accuracy. While both models achieved comparable R-squared values, the CNN's higher complexity
190 enabled it to capture more intricate patterns in the data, resulting in a lower Mean Squared Error
191 (MSE) and enhanced overall predictive capability. This highlights the CNN's capacity to excel in
192 tasks with complex relationships and spatial dependencies, showcasing its effectiveness compared to
193 the linear regression model.

194 5 Conclusion

195 In conclusion, the comparison between Convolutional Neural Networks (CNNs) and linear regression
196 underscores the trade-offs between model complexity and efficiency. While linear regression stands
197 out as a fast and sometimes effective method, suitable for simpler relationships, CNNs offer a more
198 flexible modeling approach. The latter's capacity to capture intricate patterns and spatial dependencies
199 makes it particularly well-suited for tasks where the underlying data relationships are complex.
200 Looking to the future, the potential applications of CNNs extend beyond traditional regression tasks.
201 For instance, they could be leveraged to predict and understand behavioral patterns, such as sleep
202 cycles. By integrating diverse data sources and leveraging the spatial hierarchies learned by CNNs, it
203 becomes conceivable to develop models that predict, influence, and optimize individual behaviors
204 related to mood and well-being. Such models could provide insights into how to enhance mood or
205 prevent low moods, paving the way for personalized interventions and well-being strategies based

206 on data-driven approaches. This not only demonstrates the adaptability of CNNs but also highlights
207 their potential impact in addressing complex challenges related to human behavior and mental health.